

ZIWEI ZHOU

Shanghai, China | 22307130208@m.fudan.edu.cn | Lliar-liar.github.io
github.com/Lliar-liar | [Google Scholar](#)

RESEARCH INTERESTS

Multimodal foundation models, generative AI, agent learning and self-improvement, and evaluation and benchmarking.

EDUCATION

Fudan University

B.S. in Computer Science and Technology; GPA: 3.62/4.00

Shanghai, China

Sep. 2022 – Jun. 2026

SELECTED PUBLICATIONS

- [1] **AVGen-Bench: A Task-Driven Benchmark for Multi-Granular Evaluation of Text-to-Audio-Video Generation.** [Paper] [Code]
Ziwei Zhou*, Zeyuan Lai*, Rui Wang, Yifan Yang, Zhen Xing, Yuqing Yang, Qi Dai, Lili Qiu, Chong Luo.
International Conference on Machine Learning (ICML), 2026. Poster.
- [2] **SkillOpt: Executive Strategy for Self-Evolving Agent Skills.** [Paper] [Code] [Project]
Yifan Yang*, Ziyang Gong*, Weiquan Huang*, Qihao Yang*, **Ziwei Zhou***, Zisu Huang*, *et al.*
Advances in Neural Information Processing Systems (NeurIPS), 2026. Spotlight.
14.9K+ GitHub stars and 1.3K+ forks as of Jul. 2026.
- [3] **BizGenEval: A Systematic Benchmark for Commercial Visual Content Generation.** [Paper] [Code] [Project]
Yan Li*, Zezi Zeng*, **Ziwei Zhou***, Xin Gao*, Muzhao Tian*, *et al.*
arXiv preprint, arXiv:2603.25732, 2026.
- [4] **Daily-Omni: Towards Audio-Visual Reasoning with Temporal Alignment across Modalities.** [Paper] [Code] [Project]
Ziwei Zhou, Rui Wang, Zuxuan Wu, Yu-Gang Jiang.
arXiv preprint; manuscript in preparation for IEEE Transactions on Multimedia (TMM), 2026. 60+ citations.

*Equal contribution.

RESEARCH EXPERIENCE

Microsoft Research Asia (MSRA)

Shanghai, China

Research Intern, Shanghai System and Engineering Group; Mentor: Yifan Yang

Jun. 2025 – Present

- Co-led research on multimodal generation evaluation and self-evolving agents, resulting in **AVGen-Bench** (ICML 2026), **SkillOpt** (NeurIPS 2026 Spotlight), and **BizGenEval**.
- Designed task-driven evaluation for text-to-audio-video generation across visual quality, audio quality, synchronization, text fidelity, and physical plausibility.
- Contributed to scalable data and evaluation pipelines for commercial visual generation and agent-skill optimization; released code, benchmarks, and reproducible evaluation assets.
- Implemented large-scale multimodal experimentation and distributed training pipelines with PyTorch, DeepSpeed, and Slurm.

Vision and Learning Lab (FVL), Fudan University

Shanghai, China

Research Assistant; Advisor: Prof. Zuxuan Wu

Jun. 2024 – Jun. 2026

- Led **Daily-Omni**, an audio-visual question-answering benchmark for cross-modal temporal reasoning with 684 real-world videos and 1,197 questions across six task families.
- Developed a semi-automatic annotation and verification pipeline using Gemini and DeepSeek, together with a training-free modular audio-visual reasoning baseline.
- Evaluated 24 foundation models under 37 model-modality settings; the benchmark has received **60+ citations** since its 2025 release.

TECHNICAL SKILLS

Programming: Python, C/C++, JavaScript, Bash, $\LaTeX$

Frameworks & Tools: PyTorch, TensorFlow, DeepSpeed, Docker, Slurm, Git, Hugging Face

Languages: Mandarin Chinese (native), English (TOEFL iBT: **108**; R30, L27, S23, W28)